

Sampling fidelity vs. training iteration ($h = 0.5$, exact D_{KL} to the model's own Born distribution)

Metropolis (n=20) Gibbs (n=20) Simulated Annealing (n=20) FPGA (n=20) Pegasus (+CEM) (n=20) Zephyr (+CEM) (n=20)

